

Jonathan Mandl 211399175

Danielle Hodaya Shrem 208150433

## **Part 1**

For this part, we used the same hyperparameter configuration to both our POS and NER tasks, as it delivered strong performance on each. Our neural network consists of a single hidden layer with 250 neurons. The hyperparameters were:

- **Learning rate:** 0.001
- **Epochs:** 15
- **Batch size:** 64

To handle words that appear in the development set but not in training, we added a special <UNK> token to our embedding matrix. During training, we randomly masked 15% of tokens—replacing them with <UNK>—so that the model learns a useful representation for unknown words. To handle sentence boundaries, we padded each input with <PAD> tokens on both sides (equal to the context size), and likewise included <PAD> in our embedding matrix.



